

# Class inheritance

Programação (L.EIC009)

---

José Proença (FCUP) & **João Bispo** (FEUP) – *slides by Eduardo R. B.*

- Class inheritance - fundamental notions:
  - base class and sub-class
  - inherited functionality from base class to derived class
  - protected access modifier
- Member function redefinition
  - redefinition and polymorphism
  - use of `virtual`, `override` and `final`
  - pure virtual functions puras and abstract classes
- Consolidating example
  - a class hierarchy for 2D geometric shapes

Code examples are available online at [GitHub](#).

Inheritance is normally declared as follows:

```
class SubClass : public BaseClass {  
    ...  
};
```

SubClass is called a **subclass** of the BaseClass **base class**. Other common designations: the base class is also called the **parent class** or **super-class**, the subclass is also called the **derived class** or **child class**.

SubClass inherits the functionality of BaseClass, and can define additional functionality on its own.

More advanced: SubClass may redefine the inherited functionality, and in some cases BaseClass may just define abstract functionality to be implemented by SubClass.

```
class person {  
public:  
    person(int id, const string& name);  
    person(const person& p);  
    int id() const;  
    const string& name() const;  
    ...  
private:  
    int pid;  
    string pname;  
};
```

person will be used as a base class.

```
class teacher : public person {  
public:  
    teacher(int id, const string& name,  
            const string& dept);  
    teacher(const teacher& t);  
    const string& department() const;  
    ...  
private:  
    string tdepartment;  
};
```

teacher is a subclass of person. It inherits all the definitions of person but also defines new functionality, for instance the `department()` member function.

A class may have an arbitrary number of classes, and a class hierarchy can have several levels. We could have for instance:

```
class person { ... };
```

```
// Subclasses of person
```

```
class teacher : public person { ... };
```

```
class student : public person { ... };
```

```
// Subclasses of student
```

```
class erasmus_student : public student { ... };
```

```
class working_student : public student { ... };
```

We will see a concrete example later.

Over a teacher object we can invoke the `department()` member function. Since `teacher` is a subclass of `person` we can also invoke `id()` and `name()` member functions inherited from `person`.

```
person p(123, "Joana Doa");  
std::cout << p.id() << ' '  
          << p.name() << '\n';  
  
teacher t(124, "John Doe", "Computer S  
std::cout << t.id() << ' '  
          << t.name() << ' '  
          << t.department()  
          << '\n';
```

Over a teacher object we can invoke the `department()` member function. Since teacher is a subclass of person we can also invoke `id()` and `name()` member functions inherited from person.

```
person p(123, "Joana Doa");  
std::cout << p.id() << ' '  
          << p.name() << '\\n';  
  
teacher t(124, "John Doe", "Computer Science");  
std::cout << t.id() << ' '  
          << t.name() << ' '  
          << t.department()  
          << '\\n';
```

Output:

123 Joana Doa

124 John Doe Computer Science



```
int person::id() const { return pid; }  
const std::string&  
person::name() const { return pname; }  
...  
const std::string&  
teacher::department() const { return tdepartment; }
```

`id()` and `name()` are only defined for `person`. The functionality is thus reused for `teacher`. The `teacher` subclass only needs to implement the new functionality it declares on its own such as `department()`.

Inheritance should reflect an “**is a**” relationship between the abstract notions represented by classes, e.g., a teacher is a person.

A **common design** error is to use class inheritance to model a “**has a**” relationship, for which class composition is more appropriate. For example we could have the polynomial class of previous classes formulated as:

*// BAD design*

```
class polynomial : public std::vector<fraction> { ... };
```

instead of using a `vector<fraction>` field for the coefficients. A polynomial *is not* a vector of fractions, but it *has a* (it internally employs) vector of fractions.

```
person::person
(int id, const string& name)
: pid(id), pname(name) {

}
...
teacher::teacher(int id,
const string& name,
const string& dept) :
    person(id, name), // base class constructor call
    tdepartment(department) {

}
```

**A subclass constructor must call the base class constructor.** This happens typically through a member initializer list that begins with a call to a constructor in the base class. If such a call is absent, a call to the default constructor of the base class is assumed.

```
person::person(const person &p) :  
    pid(p.pid), pname(p.pname) { }  
...  
teacher::teacher(const teacher &t)  
    : person(t), tdepartment(t.tdepartment) { }
```

Another subclass constructor example: the teacher copy constructor calls the person copy constructor.

```
class person {  
protected:  
    // directly accessible  
    // in subclasses  
    int pid;  
    std::string pname;  
    ...  
};  
  
class teacher : public person {  
    ...  
};
```

Like public and private, protected is an access modifier. **A protected declaration is accessible inside the declaring class but also any of its subclasses.** Code outside the class or subclasses cannot access protected declarations (like in the case of private)

In the above variant of person, code in the teacher subclass may directly access member fields pid and pname in person.

If we have ...

then access to pid and pname would be valid in code such as the following:

```
class person {  
protected:  
    // directly accessible  
    // in subclasses  
    int pid;  
    std::string pname;  
    ...  
};
```

```
teacher::teacher(const teacher &t)  
    : person(t),  
    tdepartment(t.tdepartment) {  
    // direct access  
    pid = pid + 1000;  
    pname = "Prof. " + pname;  
}
```

Usually inheritance is declared with a public access modifier, as in our example:

```
class teacher : public person ...
```

We will always use public in the course. But technically, it is also possible to have ...

```
class teacher : protected person ...
```

or ...

```
class teacher : private person ...
```

- public: preserves the accessibility of inherited definitions;
- protected: public declarations in the base class are changed to protected;
- private: public and protected declarations in the base class are changed to private in the subclass.

For example, suppose we have:

```
class teacher : protected person { ... };
```

Then the following code is not valid, if defined outside class teacher:

```
teacher t(123, "John Doe", "Computer Science");  
cout << t.id(); // NOT VALID!
```

**'int person::id() const' is inaccessible within this context**

id() has protected access for teacher objects, even if it remains public for person-only objects

```
person p(123, "John Doe");  
cout << p.id(); // VALID!
```



```
teacher t(...);  
person& r = t;  
person* p = &t;
```

Since `teacher` is a subclass of `person`, we can have `person` references or pointers to `teacher` objects, since the `teacher` type is a subtype of the `person` type.

The ability to use base class pointers or references to refer to objects of subclasses can be quite useful, as we will see.

On the other hand, a person reference or pointer does not necessarily refer to a teacher object. Pointer or reference assignment requires a **cast**, for example as in

```
person* p = ...;  
teacher* t = (person*) p; // cast
```

The above cast is **ok** if p really points to a teacher object. Otherwise the behavior is undefined ... bad things tend to happen ...

Excluding pointers or references, a cast cannot be used for a direct conversion from base class to subclass (we get a compilation error):

```
person p (...);  
teacher t = (teacher) p;
```

**no matching conversion for C-style cast from 'person' to 'teacher'**

What happens when we use a derived class object to construct or assign a base class object? For example, as in:

```
teacher t(124, "john doe", "Computer Science");  
person p = t;
```

Only the base class “part” of the subclass object is considered. This is known as **object slicing**.

In the example, a “slice” is taken only for the person part of t: t is treated just as a person object for the construction of p.

Both `person` and `teacher` declare the `print()` member function. The member function declaration is the same in both classes:

```
class person {  
    ...  
    void print(std::ostream& out) const;  
};  
  
class teacher : public person {  
    ...  
    void print(std::ostream& out) const;  
};
```

`teacher::print()` is in this case called a **redefinition** of `person::print()`.

`teacher::print()` may internally call `person::print()` ....

```
void person::print(std::ostream &out) const {  
    // Print id and name  
    out << "ID: " << pid << std::endl  
        << "Name: " << pname << std::endl;  
}  
  
void teacher::print(std::ostream &out) const {  
    // Call person::print  
    person::print(out);  
    // Print department info  
    out << "Department: " << tdepartment << std::endl;  
}
```

In line with the previous code of `teacher::print` then

```
teacher t(124, "John Doe", "Computer Science");  
t.print(std::cout);
```

will print

ID: 124

Name: John Doe

Department: Computer Science

Since we can have base class pointers or references what will we get in the following case?

```
teacher t(124, "John Doe", "Computer Science");  
person* pt = &t;  
pt->print(std::cout);
```

pt points to an object of type teacher but is declared as a pointer to person. Which member function is called: `person::print` or `teacher::print`?

```
teacher t(124, "John Doe", "Computer Science");  
person* pt = &t;  
pt->print(std::cout);
```

The above fragment invokes `person::print()`

ID: 124

Name: John Doe

because the declared type of `pt` refers to `person`, not `teacher`.

We can however get a different behavior if `print()` is declared to be a **virtual function** in `person`.



We can declare `print()` as a **virtual function** in `person` using the `virtual` keyword:

```
class person {  
    ...  
    virtual void print(std::ostream& out) const;  
};  
class teacher : public person {  
    ...  
    void print(std::ostream& out) const;  
};
```

In this case, `teacher::print` is called a **virtual function override**.

But what changes in practice?

If `person::print` is declared as `virtual`, then a call to `print()` through a `person` reference or pointer will use **dynamic binding**, i.e., take into account the actual **type of the object at runtime** rather than the declared type of the variable ...

Hence

```
teacher t(124, "John Doe", "Computer Science");  
person* pt = &t;  
pt->print(std::cout);
```

will call `teacher::print` ...

ID: 124

Name: John Doe

Department: Computer Science

**Good practice:** although optional, it is recommended that a virtual function override is annotated with the `override` keyword.

```
class person {  
    ...  
    virtual void print(std::ostream& out) const;  
};  
class teacher : public person {  
    ...  
    void print(std::ostream& out) const override; // use of override  
};
```

The `override` annotation explicitly signals that a function is a virtual function override. The compiler validates that the override is valid.

Example of invalid override (different 'constness'):

```
class person {  
    ...  
    virtual void print(  
        std::ostream& out) const;  
};  
class teacher : public person {  
    ...  
    // Missing const! The function is  
    // not a valid override of person::print().  
    void print(std::ostream& out) override;  
};
```

The compiler infers that the override is invalid:

```
error: non-virtual member function marked  
'override' hides virtual member function
```

```
note: hidden overloaded virtual function 'Base::print'  
      declared here: different qualifiers ('const' vs unqualified)
```

If the override annotation was missing, this would warrant at most a warning. The code would compile!

A virtual function may be declared as `final` to disallow overrides.

For `person::print` declared final

```
class person {  
    ...  
    virtual void print(std::ostream& out) const final;  
};  
  
class teacher : public person {  
    ...  
    // Not allowed!  
    void print(std::ostream& out) const override;  
};
```

We get a compilation error for the `teacher::print` override

error: declaration of 'print' overrides a 'final' function

final can also be declared for a class to disallow any subclasses ...

```
class person final { ... };  
// Subclasses not allowed  
class teacher : public person { ... };  
  
error: base 'person' is marked 'final'
```



A **pure virtual function**, sometimes also called an abstract function, is declared as `virtual ... func(...) = 0`, e.g.,

```
class person {  
    ...  
    virtual void print(std::ostream& out) const = 0;  
};
```

A class with one or more pure virtual functions is called **abstract**. An abstract class cannot be directly instantiated. **An override of the function must be defined by (concrete) sub-classes.** The function **typically also has no implementation at the base class level.**

For

```
class person {  
    ...  
    virtual void print(std::ostream& out) const = 0;  
};
```

we cannot have variables of type person ...


```
person p(123, "Joana Doa");
```

error: variable type 'person' is an abstract class

However, as teacher extends person, we can still have

```
teacher t(124, "John Doe", "Computer Science");
```

A hierarchy of classes for 2D geometrical shapes:



```
class shape {  
public:  
    virtual double area() const = 0;  
    virtual coord2d center() const = 0;  
    virtual void move(const coord2d& direction) = 0;  
    virtual ~shape();  
};
```

Since shape is abstract (it declares pure virtual functions), we cannot instantiate it directly.

```
class shape {  
public:  
    ...  
    virtual ~shape();  
}
```

**Important:** Base classes should define a virtual destructor so that subclass destructors are properly invoked when using base class pointers, e.g.

```
shape* s = new polygon ( ... );  
...  
// guarantees that destructor  
// of polygon is called  
delete s;
```

```
class ellipse : public shape {  
private:  
    coord2d ecenter; double erx; double ery;  
public:  
    ellipse(const coord2d& c, double rx, double ry) :  
        ecenter(c), erx(rx), ery(ry) { }  
    double radius_x() const { return erx; }  
    double radius_y() const { return ery; }  
    double area() const override final { return M_PI * erx * ery; }  
    coord2d center() const override final { return ecenter; }  
    void move(const coord2d& movement) override final  
        { ecenter += movement; }  
};
```

```
class ellipse : public shape {  
public:  
    double area() const override final {  
        { return M_PI * erx * ery; }  
    coord2d center() const override final  
        { return ecenter; }  
    void move(const coord2d& movement) override final  
        { ecenter += movement; }  
};
```

ellipse defines overrides for the pure virtual functions declared in shape: area(), center(), and move(). In turn, these functions are marked final disallowing their override in subclasses of ellipse (circle discussed next).

```
class circle final : public ellipse {  
public:  
    circle(const coord2d& c, double r)  
        : ellipse(c, r, r) { }  
};
```

circle is a subclass of ellipse. It cannot override any member functions of ellipse, as all of them are marked `final`.

On the other hand, circle cannot have subclasses as it is marked `final`.



(see code online)

```
class polygon : public shape { ... }  
class triangle final : public polygon { ... };  
class rectangle : public polygon { ... };  
class square final : public rectangle { ... };
```

A drawing object can group several shapes:

```
class drawing final {  
private:  
    std::vector<shape*> shapes;  
public:  
    drawing() { }  
    ~drawing() { ... }  
    void add_shape(shape* s) { ... }  
    void move_all(const coord2d& movement) { ... }  
    std::vector<shape*>& get_shapes() { ... }  
};
```

```
class drawing final {  
private:  
    std::vector<shape*> shapes;  
public:  
    ...  
    void add_shape(shape* s) {  
        shapes.push_back(s);  
    }  
};
```

To store the component shapes, drawing uses a vector of `shape*` pointers. We can not have the `vector<shape>` since `shape` is abstract. Even if `shape` was not abstract, we could not use `vector<shape>` to store objects that are subclasses of `shape` (they would be “sliced”).

Example use of drawing:

```
drawing d;  
d.add_shape(new rectangle({3, 5}, 2, 5));  
d.add_shape(new circle({1, 1}, 1));  
d.add_shape(new triangle({1, 1}, {1, 2}, {2, 1}));  
  
d.move_all({3,-1}); // move all shapes  
  
// Iterate shapes in the drawing  
for (shape *s : d.get_shapes())  
    std::cout << s->center() << ' '  
               << s->area() << std::endl;  
}
```

```
class drawing final {  
private:  
    std::vector<shape*> shapes;  
public:  
    ...  
    void move_all(const coord2d& movement) {  
        for (shape* s : shapes) {  
            s -> move(movement);  
        }  
    }  
};
```

The abstract functionality of shape lets us manipulate all (possible types of) shapes using the same interface.

drawing assumes that each contained shapes is dynamically allocated using new. Accordingly, the memory of each shape is released by the class destructor.

```
class drawing final {  
private:  
    std::vector<shape*> shapes;  
public:  
    ...  
    ~drawing() {  
        for (shape* s : shapes) {  
            delete s;  
        }  
    }  
};
```